

Group No : 43

Group Member Names:

1. SUBHRANSU MISHRA 2023AC05489
2. JAWAHARLAL RAJAN S 2023AC05504
3. SHAILESH KUMAR SINGH 2023AC05475
4. LAKSHMISRINIVAS PERAKAM 2023AC05540

1.Problem statement:

- Develop a reinforcement learning agent using dynamic programming to solve the Treasure Hunt problem in a FrozenLake environment. The agent must learn the optimal policy for navigating the lake while avoiding holes and maximizing its treasure collection.

2.Scenario:

- A treasure hunter is navigating a slippery 5x5 FrozenLake grid. The objective is to navigate through the lake collecting treasures while avoiding holes and ultimately reaching the exit (goal).

Grid positions on a 5x5 map with tiles labeled as S, F, H, G, T. The state includes the current position of the agent and whether treasures have been collected.

Objective

- The agent must learn the optimal policy π^* using dynamic programming to maximize its cumulative reward while navigating the lake.

About the environment

The environment consists of several types of tiles:

- Start (S): The initial position of the agent, safe to step.
- Frozen Tiles (F): Frozen surface, safe to step.
- Hole (H): Falling into a hole ends the game immediately (die, end).
- Goal (G): Exit point; reaching here ends the game successfully (safe, end).
- Treasure Tiles (T): Added to the environment. Stepping on these tiles awards +5 reward but does not end the game.

After stepping on a treasure tile, it becomes a frozen tile (F). The agent earns rewards as follows:

- Reaching the goal (G): +10 reward.
- Falling into a hole (H): -10 reward.
- Collecting a treasure (T): +5 reward.
- Stepping on a frozen tile (F): 0 reward.

States

- Current position of the agent (row, column).
- A boolean flag (or equivalent) for whether each treasure has been collected.

Actions

- Four possible moves: up, down, left, right

Rewards

- Goal (G): +10.
- Treasure (T): +5 per treasure.
- Hole (H): -10.
- Frozen tiles (F): 0.

Environment

Modify the FrozenLake environment in OpenAI Gym to include treasures (T) at certain positions. Inherit the original FrozenLakeEnv and modify the reset and step methods accordingly. Example grid:

S	F	F	H	T
F	H	F	F	F
F	F	F	T	F
T	F	H	F	F
F	F	F	F	G

Expected Outcomes:

1. Create the custom environment by modifying the existing "FrozenLakeNotSlippery-v0" in OpenAI Gym and Implement the dynamic programming using value iteration and policy improvement to learn the optimal policy for the Treasure Hunt problem.
2. Calculate the state-value function (V^*) for each state on the map after learning the optimal policy.
3. Compare the agent's performance with and without treasures, discussing the trade-offs in reward maximization.
4. Visualize the agent's direction on the map using the learned policy.
5. Calculate expected total reward over multiple episodes to evaluate performance.

Import required libraries and Define the custom environment - 2 Marks

```
In [1]: # Import necessary libraries
import numpy as np
import gym
from gym.envs.toy_text.frozen_lake import FrozenLakeEnv
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

Custom environment to create the given grid and respective functions that are required for the problem

Include functions to take an action, get reward, to check if episode is over

```
In [2]: class FrozenLakeTreasureEnv(FrozenLakeEnv):
        """
        Custom FrozenLake environment with treasures (T).
        Inherits from OpenAI Gym's FrozenLakeEnv.
        """
        def __init__(self, desc=None, is_slippery=False):
            """
            Initializes the environment with a custom grid.

            Args:
            - desc: Custom description of the grid (list of strings).
            - is_slippery: If True, makes the environment slippery.
            """
            if desc is None:
                raise ValueError("A custom grid (desc) must be provided for the environment.")
            super().__init__(desc=desc, map_name=None, is_slippery=is_slippery)
            self.treasure_positions = [(0, 4), (2, 3), (3, 0)] # Example treasure locations
            self.collected_treasures = set() # To track collected treasures
        def reset(self):
            """
            Resets the environment to its initial state.
            Clears the list of collected treasures.

            Returns:
            - The initial state.
            """
            self.collected_treasures = set() # Reset collected treasures
            return super().reset()
        def step(self, action):
            """
            Executes an action in the environment and returns the result.
            Adds +5 reward for collecting treasures.
            """
            # Call the parent class's step method
            result = super().step(action)

            # Unpack the returned values appropriately
            if len(result) == 5:
                next_state, reward, done, info, extra = result
                # If there's an extra value, you can ignore it or process it based on your needs
            elif len(result) == 4:
                next_state, reward, done, info = result
            elif len(result) == 3:
                next_state, reward, done = result
                info = {} # Default to an empty dictionary if `info` is not returned
            else:
                raise ValueError(f"Unexpected number of return values from step: {len(result)}")

            # Calculate the current position of the agent
            row, col = self.s // self.ncol, self.s % self.ncol # Convert state index to grid coordinates
```

```

# Check if the agent has stepped on a treasure
if (row, col) in self.treasure_positions and (row, col) not in self.collected_treasures:
    reward += 5 # Add reward for collecting a treasure
    self.collected_treasures.add((row, col)) # Mark treasure as collected

# Check if the episode ends (goal or hole)
if self.desc[row, col] in ['G', 'H']:
    done = True

return next_state, reward, done, info

def is_episode_over(self):
    """
    Checks if the episode has ended.

    Returns:
    - True if the episode has ended, False otherwise.
    """
    row, col = self.s // self.ncol, self.s % self.ncol # Current position
    return self.desc[row][col] in ['G', 'H'] # End if at Goal or Hole

def render_custom(self):
    """
    Renders the environment with additional info about collected treasures.
    """
    print("Environment Grid:")
    self.render()
    print(f"Collected Treasures: {len(self.collected_treasures)} / {len(self.treasure_positions)}")

```

Value Iteration Algorithm - 1 Mark

```

In [3]: def value_iteration(env, gamma=0.9, theta=1e-4):
    """
    Performs value iteration to compute the optimal value function (V*) and policy (π*).

    Parameters:
    - env: The environment (FrozenLakeTreasureEnv)
    - gamma: Discount factor
    - theta: Convergence threshold

    Returns:
    - V: Optimal value function for all states
    - policy: Optimal policy (actions for each state)
    """
    V = np.zeros(env.observation_space.n) # Initialize value function
    policy = np.zeros(env.observation_space.n, dtype=int) # Initialize policy

    while True:
        delta = 0 # Tracks the maximum change in the value function
        for s in range(env.observation_space.n):
            # Determine if the state is terminal
            row, col = s // env.ncol, s % env.ncol
            if env.desc[row, col] in ['H', 'G']: # Hole or Goal
                continue

            # Compute Q-values for all actions
            q_values = [
                sum(p * (r + gamma * V[s_]) for p, s_, r, done in env.P[s][a])
                for a in range(env.action_space.n)
            ]
            # Update the value function for state s
            new_value = max(q_values)
            delta = max(delta, abs(new_value - V[s]))
            V[s] = new_value
            # Update the policy to the action with the highest Q-value
            policy[s] = np.argmax(q_values)

        # Break if the value function has converged
        if delta < theta:
            break

    return V, policy

```

Policy Improvement Function - 1 Mark

```

In [4]: def policy_improvement(env, V, gamma=0.9):
    """
    Derives an improved policy based on the given value function (V).

    Parameters:
    - env: The environment (FrozenLakeTreasureEnv)
    - V: Current value function
    - gamma: Discount factor

    Returns:

```

```

- policy: Improved policy
"""
policy = np.zeros(env.observation_space.n, dtype=int)
for s in range(env.observation_space.n):
    if s in env.terminal_states: # Skip terminal states
        continue
    # Compute Q values for all actions
    q_values = [
        sum(p * (r + gamma * V[s_]) for p, s_, r, done in env.P[s][a])
        for a in range(env.action_space.n)
    ]
    policy[s] = np.argmax(q_values) # Choose the best action
return policy

```

Print the Optimal Value Function

```

In [5]: def print_value_function(V, env):
        """
        Displays the optimal value function in grid form.
        """
        grid = np.array(V).reshape(env.nrow, env.ncol)
        print("Optimal Value Function:")
        print(grid)

        custom_desc = [
            "SFFHT",
            "FHFFF",
            "FFFTF",
            "TFHFF",
            "FFFFG"
        ]

        # Assuming env is your environment and V is the optimal value function computed earlier
        env = FrozenLakeTreasureEnv(desc=custom_desc, is_slippery=False)

        # Compute the optimal value function and policy using value iteration
        V, policy = value_iteration(env)

        # Print the optimal value function
        print_value_function(V, env)

Optimal Value Function:
[[0.4782969 0.531441 0.59049 0. 0.729 0.729 0.81 0.9 1. 1.]
 [0.531441 0. 0.6561 0.729 0.81 0.9 1. 1. 1. 1.]
 [0.59049 0.6561 0.729 0.81 0.9 1. 1. 1. 1. 1.]
 [0.6561 0.729 0. 0.9 1. 1. 1. 1. 1. 1.]
 [0.729 0.81 0.9 1. 0. 0. 0. 0. 0. 0.]]

```

Visualization of the learned optimal policy - 1 Mark

```

In [6]: def visualize_policy(env, policy):
        """
        Visualizes the optimal policy as arrows on the grid.
        """
        action_symbols = ['↑', '↓', '←', '→'] # Corresponding to actions 0, 1, 2, 3
        grid = np.array(env.desc, dtype=str)
        policy_grid = grid.copy()

        for s in range(env.observation_space.n):
            row, col = s // env.ncol, s % env.ncol
            if grid[row, col] in ['H', 'G']:
                continue
            policy_grid[row, col] = action_symbols[policy[s]]

        print("Learned Optimal Policy:")
        for row in policy_grid:
            print(' '.join(row))

        custom_desc = [
            "SFFHT",
            "FHFFF",
            "FFFTF",
            "TFHFF",
            "FFFFG"
        ]

        # Updated environment initialization
        env = FrozenLakeTreasureEnv(desc=custom_desc, is_slippery=False)

        # Compute the optimal value function and policy using value iteration
        V, policy = value_iteration(env)

        # Visualize the learned optimal policy
        visualize_policy(env, policy)

```

Learned Optimal Policy:

```
↓ ← ↓ H ↓
↓ H ↓ ↓ ↓
↓ ↓ ← ↓ ↓
↓ ↓ H ↓ ↓
← ← ← ← G
```

Evaluate the policy - 1 Mark

```
In [7]: def evaluate_policy(env, policy, num_episodes=100):
        """
        Evaluates the given policy by running it over multiple episodes.

        Parameters:
        - env: The environment
        - policy: The policy to evaluate
        - num_episodes: Number of episodes to run

        Returns:
        - mean_reward: Average total reward over all episodes
        - rewards: List of rewards for each episode
        """
        rewards = []
        for _ in range(num_episodes):
            state = env.reset()
            done = False
            total_reward = 0

            while not done:
                action = policy[state]
                state, reward, done, _ = env.step(action)
                total_reward += reward

            rewards.append(total_reward)
        mean_reward = np.mean(rewards)
        return mean_reward, rewards
```

Main Execution

```
In [8]: if __name__ == "__main__":
        # Define a custom 5x5 grid
        custom_desc = [
            "SFFHT",
            "FHFFF",
            "FFFTF",
            "TFHFF",
            "FFFFG"
        ]

        # Initialize the custom environment with the custom grid
        env = FrozenLakeTreasureEnv(desc=custom_desc, is_slippery=False)

        # Perform value iteration
        V, policy = value_iteration(env)

        # Print the optimal value function
        print_value_function(V, env)

        # Visualize the learned policy
        visualize_policy(env, policy)

        # Evaluate the policy
        mean_reward, rewards = evaluate_policy(env, policy)
        print(f"Mean Reward over {len(rewards)} episodes: {mean_reward}")
```

Optimal Value Function:

```
[[0.4782969 0.531441 0.59049 0. 0.729 ]
 [0.531441 0. 0.6561 0.729 0.81 ]
 [0.59049 0.6561 0.729 0.81 0.9 ]
 [0.6561 0.729 0. 0.9 1. ]
 [0.729 0.81 0.9 1. 0. ]]
```

Learned Optimal Policy:

```
↓ ← ↓ H ↓
↓ H ↓ ↓ ↓
↓ ↓ ← ↓ ↓
↓ ↓ H ↓ ↓
← ← ← ← G
```

Mean Reward over 100 episodes: 6.0